

Phonebook Lookup — Group Report (Parts A, B, D)

Group: G4 — Searching & Hash tables (§2.2) **Members:** **Topic:** Linear · Binary · Interpolation search · hashing & collisions **Demo:** Phone-book lookup with a hash table ($n \geq 100\,000$)

Written form of Parts A, B and D of the assignment
(C0B_Group_Project_Guide.pdf). Build/run instructions and the machine
used are in [../README.md](#).

Part A — Position & purpose

A.1 Syllabus map

Place each structure on the syllabus outline: chapter, section, page, and which learning outcome (LO1 / LO2 / LO3) it serves. References are to the course syllabus (3-Data Structure and Algorithm-CLC.pdf, INT1306_CLC).

Structure	Syllabus ref	LO	Built on	Built on top of it
Linear search	§2.2.2, p.3	LO2	array (unsorted)	—
Binary search	§2.2.3, p.3	LO2	sorted array	—
Interpolation search	§2.2.4, p.4	LO2	sorted array, uniform keys	—
Hashing & collisions	§2.2.5, p.4	LO2 + LO3	array (buckets) + linked list (chaining)	—

A.2 Motivation — one paragraph per structure

For each structure answer explicitly:

- What problem does it exist to solve? What was unacceptable before it?
- Which structures does it build on, and which are built on top of it?
- When is it the **wrong** choice? Name a concrete case where you would not use it.

A.3 Sources

- Course syllabus — *Data Structures & Algorithms (INT1306_CLC)*, PTIT,

3-Data Structure and Algorithm-CLC.pdf, accessed 2026-09-11. *

Part B — Algorithms & complexity

B.1 Operation table

Every operation in the topic is tabulated. The demo implements linear search, binary search (on a maintained sorted phone index), and the hash table; interpolation search is part of the topic (§2.2.4) and is included for completeness but not in the demo (it needs uniformly distributed keys; phone strings are not). A number without a justification scores nothing.

Operation	Best	Average	Worst	Space	Why
Linear search by phone	$O(1)$	$O(n)$	$O(n)$	$O(1)$	No index — every element is compared until found; first match is $O(1)$, a miss scans all n .
Binary search by phone	$O(1)$	$O(\log n)$	$O(\log n)$	$O(1)$	Halves the search space on each probe; requires a sorted array.
Interpolation search	$O(1)$	$O(\log \log n)$	$O(n)$	$O(1)$	Probes by value proportion; $O(\log \log n)$ only for uniformly distributed keys, degenerates to $O(n)$ on skewed data.
Hash insert	$O(1)$	$O(1)$	$O(n)$	$O(1)$	Hash to a bucket and append to its chain; expected $O(1)$ at constant load factor, $O(n)$ if every key collides into one chain.
Hash search	$O(1)$	$O(1)$	$O(n)$	$O(1)$	Hash to a bucket and walk the chain; expected $O(1)$, worst case a single long chain.
Hash delete	$O(1)$	$O(1)$	$O(n)$	$O(1)$	Find the node in the chain, then unlink it; same bounds as search.
Rehash (load > 0.75)	$O(n)$	$O(n)$	$O(n)$	$O(n)$	Every key is reinserted into a new, doubled bucket array; amortised $O(1)$ per insert.

B.2 Implementation notes (from the C++ source)

Moved verbatim from the README — the concrete choices behind the bounds above.

- CSV (src/phonebook.cpp:30-83): quoted "Do, Thanh Tuan" with ""

escapes, unquoted split on **last** comma, trim spaces, strip wrapping quotes on phone, drop trailing `\r`, skip blank lines, skip rows with empty name or non-digit phone.

- **Insert** (`src/phonebook.cpp:143-164`): reject empty, reject

non-digit phone (`Invalid phone number`), reject duplicate phone via hash (`Phone number is already exist`), else store `CapitalizeFirst(ToLower(name))` (ASCII-only) + index in hash table.

- **Delete** (`deleteContactByPhone`, option 8): $O(n)$ — vector erase

shifts the tail + full hash rebuild so order is preserved. Missing phone prints `Phone number not found`.

- **Search**: `searchLinearByPhone` exact scan, `searchHashByPhone` chained

lookup, `searchBinaryByPhone` hand-written binary search on the sorted phone index ($O(\log n)$, resolves to the contact via the hash table), `searchLinearByName` case-insensitive scan via `toLowerCase`.

- **Sorted index** (`buildSortedIndex`): bulk load pushes contacts + hash

entries directly, then sorts the phone copy once ($O(n \log n)$); each `insertContact` keeps it sorted via hand `lowerBound` ($O(\log n)$) + shift ($O(n)$) and prints the split hash `insert / sorted-index insert` timing; `deleteContactByPhone` rebuilds it (delete is $O(n)$ anyway).

- **Hash** (`src/hashtable.cpp`): `hash = hash*31 + (unsigned char)(c-'0')`

with 64-bit wrap, `% numBuckets`; start 101 $\rightarrow$ `nextPrime`; rehash at load > 0.75 to `nextPrime(2x)`; `isPrime` uses `6k±1`.

- **Timer** (`include/timer.hpp`): `timeIt ms`, `benchmark(work, 5)`

best-of-5, `printTaskDuration` prints `\nTook: Xms`.

- **Benchmark** (`src/main.cpp:11-63`, option 0): auto-loads CSV if empty,

picks `first / middle / random / last / miss` phones (`miss` is `0000000000`, not in dataset — true worst case for hash/binary), times linear vs hash vs binary (sorted index) 5 runs each.

Part D — Cross-language comparison

D.1 What is being compared

Same handwritten algorithm, same input, same machine: a hand-rolled chained hash table, a hand-rolled linear scan, and a hand-rolled binary search on a maintained sorted phone index over `contacts_100k.csv` ($n=100000$) in all five languages — no `dict`, `map`, or `HashMap` anywhere. This is a *language + runtime* comparison of equivalent code. It is deliberately **not** hand-written C++ vs Python's built-in `dict`: that would compare a teaching implementation against a production hash table (open addressing, optimized C) and would answer a different question.

D.2 The same program in five languages

Same prompts, same messages, same file split as C++. Differences are only what the runtime forces (64-bit wrap, `stdin`, timing API).

C++	Python	Go	JavaScript (Node)	Java
<code>src/main.cpp</code>	<code>python/phonebook/main.py</code>	<code>go/phonebook/main.go</code>	<code>javascript/phonebook/src/main.js</code>	<code>java/phonebook/src/com/phonebook</code>
<code>src/phonebook.cpp</code> + <code>include/phonebook.hpp</code>	<code>phonebook.py</code>	<code>phonebook.go</code>	<code>src/phonebook.js</code>	<code>PhoneBook.java</code>
<code>src/hashtable.cpp</code> + <code>include/hashtable.hpp</code>	<code>hashtable.py</code>	<code>hashtable.go</code>	<code>src/hashtable.js</code>	<code>HashTable.java</code>
<code>include/timer.hpp</code>	<code>timer.py</code>	<code>timer.go</code>	<code>src/timer.js</code>	<code>Timer.java</code>
<code>include/contact.hpp</code>	<code>contact.py</code>	<code>contact.go</code>	<code>src/contact.js</code>	<code>Contact.java</code>


D.3 Comparison table

Fill every cell. The last row — where you would use it — is what earns the marks.

Criterion	C++	Python	Java	JavaScript	Go
Built-in structure used	—	—	—	—	—
Manual implementation effort					
Runtime on your dataset					
Memory footprint					
Memory management					
Type safety					
Readability of the code					
Where you would use it					


D.4 Runtime growth — line chart

Best-of-5 wall-clock ms at the `last` position (index $n-1$, the same position in every language) across $n = 50, 10k, 100k, 200k, 500k, 1M$ — the runtime growth curves of all 5 languages overlaid on one graph (log-log). Linear search grows $\sim$ linearly with n ; hash lookup stays flat ($\sim O(1)$) and binary search stays nearly flat ($O(\log n)$) — both position-independent, unlike linear.




Per-language view — one chart per language, every algo (linear / hash / binary) on every case (first / middle / random / last / miss) vs n , log-log. Color = algo, linestyle = case (first solid, middle dashdot, random dashed, last dotted, miss solid). Each PNG answers "how fast is *this* language on every algorithm" without cross-language clutter:


cpp: every algo on every case vs n




python: every algo on every case vs n




go: every algo on every case vs n




js: every algo on every case vs n






Per-language best / avg / worst per algo — one chart per language, x=algo, 3 bars best / avg / worst at the largest n (log y). best=min(hit), avg=mean(hit), worst=max(hit,miss) where hit=first/middle/random/last. This is the "best worst and avg case on each algo" view: linear fans out (best<<avg<<worst), hash/binary stay flat.


cpp: best / avg / worst per algo (n=1000000)




python: best / avg / worst per algo (n=1000000)




go: best / avg / worst per algo (n=1000000)



js: best / avg / worst per algo (n=1000000)





D.5 Memory profiling — bar chart

Guide: use language-specific profilers to capture exact peak memory: Valgrind/Massif (C++), tracemalloc (Python), VisualVM (Java), Chrome DevTools (Node/JS), pprof (Go).

D.6 Statistical fairness

Guide: run each test at least 5 times (warm-up runs do not count for Java/JS JIT compilers); report the average and the variance; explain why Java is slow on run 1 but fast on run 5 (JIT compilation).

D.7 Measurement method

Option 0 now covers 5 targets (first/middle/random/last/miss) and prints best-of-5 per target:

```
[role index N phone P]
Linear best of 5: Xms. (index N)
Hash best of 5: Yms. (index N, position-independent)
Binary best of 5: Xms. (index N, sorted index, position-independent)
```

miss uses phone 0000000000 (not in dataset) — true worst case for hash (full chain) and binary (log n probes); for linear it is also worst (full scan). middle is n/2 — linear avg and binary best (1 probe).

Run the full matrix:

```
make run-benchmark # all 5 langs, every run → benchmark/results.csv (overwrite)
```


Each program also supports batch mode directly (75 rows each — first/middle/random/last/miss × linear/hash/binary × 5 runs; compile the C++ demo first):

```
make # builds ./build/cpp/demo; other languages need no build step
./build/cpp/demo --benchmark-csv benchmark/results.csv data/contacts_100k.csv
python3 python/phonebook/main.py --benchmark-csv benchmark/results.csv --append data/contacts_100k.csv
(cd go/phonebook && go run . --benchmark-csv ../../benchmark/results.csv --append ../../data/contacts_100k.csv)
(cd javascript/phonebook && node src/main.js --benchmark-csv ../../benchmark/results.csv --append ../../data/contacts_100k.csv)
(cd java/phonebook && javac -d out src/com/phonebook/*.java && java -cp out com.phonebook.Main --benchmark-csv ../../benchmark/results.csv --append ../../data/contacts_100k.csv)
```

benchmark/results.csv columns: language,dataset,n,case,algo,run,ms,timestamp,toolchain,target_index,phone — 375 rows single-size (5 langs × 5 cases × 3 algos × 5 runs), 2250 rows multi-size (× 6 sizes). Batch mode seeds RNG with 42 so the random target is reproducible within a language (each language's RNG differs, so targets differ across languages — target_index,phone columns make that auditable). Option 0 stays unseeded. miss has target_index=-1.


make run-benchmark also draws benchmark/plot.png at the end via benchmark/plot.py (three linear-scale panels — linear scan | hash lookup | binary search — best of 5 per cell with value labels; split axes so bar heights stay in true ratio and the μs hash/binary bars aren't flattened by the ms linear bars;

needs matplotlib — without it the plot step prints `plot skipped` and the CSV is still produced). The script parses every language's `ms` float format (C++ `setprecision(17)`, Python `repr`, Go `%g`, JS double, Java `Double.toString`) defensively: bad rows are skipped with a warning, never a crash. `make run-benchmark-sizes` additionally draws `plot-runtime-vs-n.png` (line chart), `plot-unified-*.png` (every algo on every case per language), and `plot-per-algo-*.png` (best/avg/worst per algo).




▼ Per-group charts (one image per case × algo, horizontal, fastest first)


first-linear across languages, fastest first (data/contacts_50.csv, n=50)




first-hash across languages, fastest first (data/contacts_50.csv, n=50)




first-binary across languages, fastest first (data/contacts_50.csv, n=50)




middle-linear across languages, fastest first (data/contacts_50.csv, n=50)




middle-hash across languages, fastest first (data/contacts_50.csv, n=50)




middle-binary across languages, fastest first (data/contacts_50.csv, n=50)




random-linear across languages, fastest first (data/contacts_50.csv, n=50)




random-hash across languages, fastest first (data/contacts_50.csv, n=50)




random-binary across languages, fastest first (data/contacts_50.csv, n=50)




last-linear across languages, fastest first (data/contacts_50.csv, n=50)




last-hash across languages, fastest first (data/contacts_50.csv, n=50)




last-binary across languages, fastest first (data/contacts_50.csv, n=50)




miss-linear across languages, fastest first (data/contacts_50.csv, n=50)



miss-hash across languages, fastest first (data/contacts_50.csv, n=50)





D.8 Explaining the gap, not just observing it

"Python is slower" is an observation; the mechanism is the answer:

- **C++ (-O2)** is the baseline: `vector<Contact>` stores structs

contiguously, short strings sit inside the object (SSO), and the scan compiles to a tight machine-code loop with no per-element overhead.

- **Python** pays per element: every `Contact`, `str`, and `int` is a

heap-boxed `PyObject` reached through a pointer, and each loop iteration runs CPython bytecode dispatch plus attribute-dict lookups. That is why `last-linear` is ~10× the others — 100k interpreter steps vs 100k machine instructions. The hash bars stay small because $O(1)$ steps mean the per-op tax applies ~once, not 100k times.

- **Go** is compiled with value-type slices, so its scan is a real

machine loop (bounds checks are nearly free here) — expect it closest to C++. Its gap is runtime checks + GC write barriers, not interpretation.

- **JavaScript (V8)** stores an array of heap objects (pointer chasing +

hidden-class checks per access) and must JIT-warm: best-of-5 lets the optimizing compiler settle, but run-to-run spread is JIT tier-up, not noise in your code.

- **Java** is the same story on the JVM: `ArrayList<Contact>` holds

references, the C2 compiler kicks in after ~10k loop iterations (one 100k scan trips it), and G1 may pause — startup is *not* measured, only the timed search, which flatters Java fairly since all languages get the same treatment.

Same position, except random. `first` (index 0) and `last` (index $n-1$) hit the identical phone in every language, so those bars are directly comparable. The `random` target is seeded (42) but each language has its own RNG, so the index differs per language — check `target_index` in the CSV before comparing `random-*` bars across languages.

What you may claim. The portable, algorithmic claim is the *trend within each language*: linear grows first → last, hash/binary stay ~constant — it holds in all five, which is the DSA point. Binary costs $O(\log n)$ probes against a sorted phone copy (built once at load, kept sorted on insert), so it tracks hash closely while linear diverges. Cross-language magnitude differences are runtime properties; cite them only with the mechanism above, plus your machine spec (CPU/RAM, OS, toolchains, commit).

D.9 Recommendation

D.10 Measurement tables

Best-of-5 wall-clock ms per target (`first` / `random` / `last`) from a fresh 100k run on this machine (uncommitted binary-port code; re-run `make run-benchmark` after committing to refresh `benchmark/results.csv`).

Run	Dataset	CPU/RAM, OS	Toolchain	Commit	n	target	Linear runs (ms)	Hash runs (ms)	Binary runs (ms)
C++ ex	contacts_100k	i5-1135G7, 15 GiB, Omarchy	g++ 16.2.1, -O2	uncommitted	100000	first/random/last	0.00002 / 0.089025 / 0.365182	0.000037 / 0.000037 / 0.000042	0.000111 / 0.000125 / 0.000189
py ex	contacts_100k	i5-1135G7, 15 GiB, Omarchy	Python 3.14.7	uncommitted	100000	first/random/last	0.000225 / 4.03942 / 5.07507	0.001378 / 0.001283 / 0.001919	0.002869 / 0.002241 / 0.003748
go ex	contacts_100k	i5-1135G7, 15 GiB, Omarchy	go1.27.1	uncommitted	100000	first/random/last	0.000031 / 0.323417 / 0.441653	0.00005 / 0.000039 / 0.000046	0.000221 / 0.00023 / 0.000242
js ex	contacts_100k	i5-1135G7, 15 GiB, Omarchy	node v26.7.0	uncommitted	100000	first/random/last	0.000311 / 0.469038 / 0.933914	0.000938 / 0.00092 / 0.000743	0.001878 / 0.001655 / 0.00149
java ex	contacts_100k	i5-1135G7, 15 GiB, Omarchy	javac 27	uncommitted	100000	first/random/last	0.000659 / 0.698453 / 1.71081	0.000497 / 0.000475 / 0.000566	0.002898 / 0.001821 / 0.002278

D.11 Test conditions (machine & toolchain)

A measurement without its conditions is not a measurement. Same table as README.md — fill once, keep both in sync.

Component	Value
CPU	Intel Core i5-1135G7 @ 2.40 GHz (4 cores / 8 threads)
RAM	15 GiB
OS	Omarchy (Arch Linux), kernel 7.2.3-arch1-3
C++	g++ 16.2.1 (GCC), -O2 (see Makefile)
Python	Python 3.14.7
Go	go1.27.1 linux/amd64
Node	v26.7.0
Java	javac 27 (JDK 27)
Dataset	data/contacts_100k.csv (n=100000), seed 42
Commit	adc54407393a2d8b82260a8af5695ad077df7641

Appendix

A.1 Test conditions

See Part D.11 — CPU/RAM, OS, compiler versions, optimisation flags, dataset and how it was generated.

A.2 Dataset

See data/README.md. Name,Phone per line, no header. Regenerate with python3 data/generate.py --seed 42. Sizes: 50 / 10k / 100k / 200k / 500k / 1m — nested prefixes of one seeded stream, so scaling results are comparable. Use contacts_50.csv for demos.

A.3 Port implementation notes

- **Python:** time.perf_counter, random.randint, 64-bit wrap via & 0xFFFFFFFFFFFFFFFF, (ord(c)-48) & 0xFF to match (unsigned char)(c-'0').
- **Go:** uint64 wraps like C++ size_t, time.Now, math/rand, bufio.Reader line input.
- **JavaScript (Node):** perf_hooks.performance.now, Math.random, readline/promises async CLI, BigInt with & 0xFFFFFFFFFFFFFFFFn for hash wrap, fs.readFileSync.
- **Java:** System.nanoTime()/1e6, java.util.Random, Scanner hasNextInt/nextLine mirroring cin >> choice, long wrap + Long.remainderUnsigned for hash mod. out/ is git-ignored.

A.4 Sources

A.5 Contribution statement

A.6 AI-use declaration